


Final Schedule

	course	time_slot	room	facilitator
0	SLA101B	1000	BEACH 301	GLEN
1	SLA449	1000	LOFT 310	TYLER
2	SLA191B	1100	FRANK 119	GLEN
3	SLA291	1100	BEACH 301	LOCK
4	SLA451	1100	SLATER 003	TYLER
5	SLA191A	1300	BEACH 201	GLEN
6	SLA304	1300	SLATER 003	GLEN
7	SLA201	1400	LOFT 206	GLEN
8	SLA101A	1500	BEACH 301	LOCK
9	SLA303	1500	BEACH 201	GLEN
10	SLA394	1500	FRANK 119	TYLER

Fitness over the generations



i) What were the challenges in writing the program? Or did it seem to go smoothly from the beginning?

Creating the program to create chromosomes, the initial generation and subsequent generations went very easily. However, I tried to implement parallelism with a new library, `asyncio`, and it didn't work with jupyter notebooks. I've used a parallel backend for others, but I did not want to take the time to refactor the program considering 500 generations took about 20 minutes. However, I left the `asyncio` in place for use in local environments. Once I had the code running, I thought things were going great, but when I printed the schedule, it had assigned one instructor to every class. I was certain something was wrong with my function that evaluated fitness, but I eventually figured out that my interpretation of,

*"If any facilitator is scheduled for consecutive time slots,
same rules for SLA101 and SLA191 below."*

resulted in a reward of .5 for every consecutive time slot that outweighed the penalty for simultaneous classes and over 4 classes. I elected to only award the consecutive time slots for classes SLA101 and SLA191.

ii) What do you think of the schedule your program produced? Does it have anything that still looks odd or out of place?

I find the schedule my program produced acceptable. However, it would be better if it placed the 101 and 191 classes farther apart, but the reward for subsequent facilitators and penalty for over 4 classes makes this harder to evolve. I tried playing around with generation size, mutation rate, the mutation decay rate, and how many survivors there were between generations, but it never quite created a schedule with 101/191 in the first two hours and last two hours.

iii) How would you improve the program, or change the fitness function?

A blind genetic algorithm has uses, but I couldn't help but think that adding constraints would greatly improve efficiency. Simply not allowing the program to double book rooms or facilitators would allow the reward to focus on optimality. I considered just killing population members that violated this, but they are bred out relatively quickly. However, if children in each generation weren't allowed to break constraints, it would vastly improve the gene pool. I also felt the room constraints were too lenient.

iv) Anything else you feel like discussing, asking about, bragging about, etc.

I was satisfied with the intuitiveness of creating a genetic algorithm. My partner, a biology major, isn't normally that intrigued by computer science, but she thought this was really cool. It was interesting to see how conflicting rewards and punishments interplayed, even if it took me hours to figure out the algorithm was exploiting the heuristic against my intent. It was good to play with `asyncio`, but, unless I see a clear need, I'll probably stick with the parallelism libraries I'm familiar with.